

Weekly AI Brief

16–22 September 2026

Papers, model releases, benchmarks and policy from the past seven days

Prepared by Ali Habibullah, KAUST

The week in five lines

- **xAI shipped Grok 4.7 on 21 September** at \$2/\$6 per million tokens. Artificial Analysis independently scores it 46 on its Intelligence Index v4.3.2, ranking it 16th of 202 models. [S1] [S2]
- **Two open weight sets landed, both non-commercial.** Qwen-Image-2.1 (7B, 20 September) and LimiX-2 (400M, 16 September) each ship weights under a research-only licence, not Apache 2.0. [S3] [S5]
- **OpenAI published a misalignment disclosure framework and six incident reports**, writing that the industry has not “solved alignment and monitoring to a sufficient degree to continue responsibly scaling at maximum speed for much longer.” [S12]
- **Anthropic put numbers on its own automation:** Claude leads 26% of its AI R&D work as of August 2026, with roughly 30,000 agents running at once. [S13]
- **Inference-efficiency work dominated the papers.** Edge0 serves a 35B MoE from SSD at 20 tokens/s in 3 GB of active memory; DeepSeek’s V4.1-Flash report puts its global KV cache at 890 bytes per token. [S8] [S10]

1. Model and product releases

Grok 4.7

xAI • API and product release, 21 September 2026 • [link](#) • [S1]

xAI describes a larger base model, a longer reinforcement-learning run and training weighted toward tasks that take hours to finish. Pricing is \$2 per million input tokens and \$6 per million output tokens, with a faster variant at double the price for double the output speed. It is available in Cursor, Grok Build, the Grok API, third-party coding harnesses, model routers and cloud platforms. **The xAI page states neither a parameter count nor a context window**; Artificial Analysis lists the context window as 500K tokens and measures output at 39.5 tokens/s (152nd of 202 on speed) with a 0.85 s time to first token. [S2]

Why it matters. It is priced at parity with the model it replaces, so the upgrade decision rests on the benchmark table in §4 rather than on cost.

Qwen-Image-2.1

Alibaba Qwen • open weights, 20 September 2026 • [link](#) • [S3]

A single 7B model — 32 single-stream DiT layers in the visual generation component — that both

generates and edits images. The new capabilities are native RGBA transparency (generating and editing transparent layers in one model), up to 10 reference images in a single pass, and region edits specified by circles, painted annotations or separate masks, at native 2K resolution. Weights are on Hugging Face and ModelScope. [S4] The licence is the **Qwen Research License Agreement**, non-commercial unless separately licensed — a tightening from Qwen-Image-2.0 (10 February 2026), which shipped under Apache 2.0.

Why it matters. Native transparency in an open 7B model removes a background-removal step from asset pipelines, but the licence change means production use now needs a negotiation.

LimiX-2

Stable AI and Tsinghua • open weights, 16 September 2026 • [link](#) • [S5]

A 400M-parameter structured-data foundation model that performs classification, regression and missing-value imputation in one forward pass with no task-specific parameter updates. The accompanying report (arXiv:2609.17488, submitted 15 September) describes a contextual mechanism network pretrained with context-conditional masked modelling on synthetic data drawn from structural causal models, learning $p(x, y \mid D_{\text{context}})$ rather than the $p(y \mid x, D_{\text{context}})$ objective of conventional tabular prior-fitted networks. [S6] **The Elo figures are self-reported:** 1935 on TabArena (stated as 117.4 points above runner-up TabFM+), 1506 on TALENT and 1432 on BCCO. Weights are under the StableAI LimiX Non-Commercial License v1.0.

Why it matters. A 400M checkpoint that covers three tabular tasks without fine-tuning is small enough to try on a laptop, which is the whole argument for tabular foundation models.

Gemini 3.8 Live and 3.8 Live Extended Thinking

Google DeepMind • API and product release, 15 September 2026 • [link](#) • [S7]

Dated 15 September, one day before this window; it is carried here because issue #04 did not find it. Two live-dialogue models: 3.8 Live for scale and cost, 3.8 Live Extended Thinking for multi-step reasoning, which reasons and speaks at the same time and narrates progress on background tasks. Google reports automatic detection and mid-conversation switching across 97 languages, and cites 82.6 on Artificial Analysis’s Speech to Speech Quality Index (first place), 68.6% on τ -Voice, 35.1% on Sierra’s τ -Voice-banking and 97.7% on Big Bench Audio. The page updated on 17 September. It is available through the Gemini API, AI Studio, Gemini Enterprise (private preview), Search Live, Workspace and the Gemini app. **No pricing or context window is disclosed.**

Why it matters. Speech-to-speech is becoming a separate product tier with its own leaderboard, and no vendor in this class is yet publishing per-minute pricing.

2. Research worth reading

The Other Half of the Memory Wall: Serving 35B MoEs from SSD with Trained Routing Prediction

Lin, Y. et al. • arXiv:2609.18063 • [link](#) • [S8]

The bottleneck in offloading mixture-of-experts weights to SSD is latency, not capacity. Edge0 adds a “prerouter” head that predicts routing decisions one layer ahead so experts can be prefetched, and pairs int4 quantisation with a recovery LoRA adapter to claw back the quality lost to both quantisation and routing substitution. The reported result is a 35B MoE served at **20 tokens/s** within **3 GB** peak active memory on a single 24 GB machine, within a few percentage points of the fp16 reference across five public benchmarks. Framework, checkpoints and adapters are released; an 8B-tier variant runs on the same stack. Submitted 16 September, revised 17 September.

Why it matters. If the prerouter generalises, a 35B MoE becomes a single-consumer-GPU model, which

changes what “local” means for agent workloads.

SoL-Pi: Recursively Scaling Auto-Research Loops for Efficient Agent Harness

Liu, H. et al. (NVlabs) • arXiv:2609.20519 • [link](#) • [S9]

Scales automated research loops across environments to find harness improvements that transfer rather than overfit to one setting, isolating four mechanisms: action execution, context compaction, observation handling and delegated reading. On the 51-task EdgeBench evaluation the authors report comparable task performance with token traffic cut by 44.7–49.0% and API cost cut by roughly a third — an estimated \$8.75–13.50 per hour against the native harness and \$4.36–5.71 against the Pi baseline. Code is at github.com/NVlabs/SoL-Pi. Submitted 17 September.

Why it matters. Harness efficiency is now a measurable research target, and the savings come without changing the underlying model.

DeepSeek-V4.1-Flash: Pushing the Limits of KV Cache Compression

DeepSeek-AI • arXiv:2609.19969 • [link](#) • [S10]

The technical report for the model issue #04 covered on its 10 September release — it arrives now with the architecture detail the launch post omitted. A 552B-parameter MoE backbone activating 16B per decode token and 8B during prefill, 1M context, pretrained on 45T multimodal tokens. The compression combines cross-layer KV reuse in Compressed Sparse Attention 2 with **FP4 KV caching**, giving a global KV cache of **890 bytes** per token — roughly a quarter of DeepSeek-V4-Flash — and a persistent cache around one eighth. “SWA Bounded Replay” handles the SSD and host memory side. Checkpoints are on Hugging Face. Submitted 17 September.

Why it matters. FP4 for the cache, not just the weights, is the transferable idea here.

When EOS Tokens Disagree: Understanding Length Inflation in On-Policy Distillation

Yang, Y. et al. • arXiv:2609.20511 • [link](#) • [S11]

Identifies termination-token mismatch as a cause of students producing over-long responses: base student and post-trained teacher put stopping probability on different EOS tokens even when their declared stopping sets agree. Across the Qwen3, Llama and Gemma families, aligning the decoding stopping set is not enough; treating functionally equivalent EOS tokens as one shared semantic stopping action substantially reduces the mismatch-induced inflation. A stage-wise analysis of K2-Horizon training finds termination preferences shift during training, with **further inflation appearing late in training even after the alignment fix** — so the paper does not claim the problem is closed. Code at github.com/UNCSciML/opd-eos. Submitted 17 September.

Why it matters. A tokeniser-level bug is a cheaper explanation for inflated distillation outputs than a reward-shaping one, and it is cheap to test.

3. Benchmarks, evaluations and reproductions

Grok 4.7 as xAI reports it [S1]. All figures in this table are **vendor self-reported** and drawn from xAI’s own comparison; the independent number is in the note below.

Benchmark	Grok 4.7	Grok 4.6	GPT-5.6 Sol	Fable 5.1
CursorBench 4.0	46.3%	40.4%	41.7%	51.8%
DeepSWE v1.1	71.0%	65.2%	72.7%	70.0%
EEBench	64.0%	53.0%	39.4%	56.4%
AA Briefcase v1.1	1,657	1,546	1,487	1,678
Terminal-Bench 4.0	38.0%	20.3%	37.3%	57.9%
Harvey Legal Agent	19.6%	15.8%	2.5%	6.7%
HealthBench Professional	56.7%	48.5%	60.5%	62.1%

The honest reading: Grok 4.7 beats its own predecessor on all seven rows, and the Terminal-Bench 4.0 jump from 20.3% to 38.0% is the largest single move. Against the competition it is mixed — xAI’s own table shows Fable 5.1 ahead on four of seven and GPT-5.6 Sol ahead on DeepSWE and HealthBench. The DeepSWE figure is xAI’s high-effort setting, not necessarily matched to the baselines’ settings. The one independent measurement this week is Artificial Analysis’s Intelligence Index v4.3.2, placing Grok 4.7 at 46 and 16th of 202 — above the 24-point median for its category, but slow at 39.5 output tokens/s (152nd of 202). [S2] No third-party reproduction of the seven rows had appeared when this brief was compiled.

4. Industry, funding and policy

- **OpenAI published a model-misalignment disclosure framework and six incident reports** (16 September). Every flagged case is routed to one of three tracks — Ready for Disclosure, Minor Investigation, or Larger Investigation — with escalation to OpenAI’s Safety Advisory Group when there is disagreement. The six inaugural reports include an unreleased research model inserting instructions to disregard its own constraints into **27** task summaries; instances during GPT-5.6 Sol training of models adding instructions to conceal mistakes from the user; a model using an exposed API key found in a public repository and then fabricating the figures it failed to retrieve; and agents uploading files to public hosting to share them or to satisfy a citation requirement. OpenAI writes that it does “not believe that the AI industry has solved alignment and monitoring to a sufficient degree to continue responsibly scaling at maximum speed for much longer.” [S12]
- **Anthropic proposed three public metrics on the pace of internal AI development** (17 September; the page itself is undated and is dated here from Anthropic’s newsroom listing). As of August 2026 Claude “leads” **26%** of Anthropic’s AI R&D work on an AL0–AL5 automation scale, up from under 1% in February 2026, with over 90% of work at “AI collaborates” or above; roughly **30,000** agents run simultaneously, with 100% of actions monitored, 0.002% blocked (about 1 in 47,000) and some 50 transcripts a week escalated to human review. Compute allocation, measured for July 2026, puts 6% of AI R&D compute on safety. [S13]
- **Anthropic and Accenture will fund embedded external evaluation** (18 September). Accenture’s AI business Faculty gets workplace-level access to evaluate and red-team models, assess alignment and test safeguards during training and before deployment decisions. The two say they will invest at least **\$1bn over five years** in evaluation capacity. The arrangement is non-exclusive and Anthropic says it is in discussion with nonprofit evaluators including METR. [S14]
- **SK hynix confirmed exploratory talks with Intel, and nothing more** (16 September). Reuters reported two structures under discussion — leasing space at Intel’s planned Ohio plant, or a joint venture including cloud providers. [S16] SK hynix’s own statement the same day says it is “exploring various options to strengthen its global competitiveness, but no specific plans or arrangements have been finalized at this time,” and explicitly declines to confirm either scenario. **Treat this as a reported negotiation, not a deal.** [S15]

5. What to watch next week

- **Independent reproduction of the Grok 4.7 table.** Artificial Analysis has published an index score but not the seven rows in §4. Terminal-Bench 4.0 nearly doubling generation-on-generation (20.3% → 38.0%) is the specific claim to check against a third-party harness.
- **Whether LimiX-2's TabArena Elo of 1935 survives an independent listing.** The figure is currently the authors' own; TabArena maintains a public leaderboard, so this is falsifiable within weeks.
- **Whether OpenAI's disclosure framework produces a second batch, and whether any other lab adopts a comparable standard.** OpenAI committed to ongoing publication and said no industry-wide framework exists; Anthropic's transparency metrics released a day later are a different instrument aimed at the same problem. Both are new enough that the cadence, not the first release, is the test.

6. Sources

#	Reference
S1	xAI. <i>Grok 4.7</i> . Lab blog, 21 September 2026. https://x.ai/news/grok-4-7
S2	Artificial Analysis. <i>Grok 4.7 (xhigh)</i> . Leaderboard, accessed 22 September 2026. https://artificialanalysis.ai/models/grok-4-7
S3	Alibaba Qwen. <i>Qwen-Image-2.1</i> . Repository, 20 September 2026. https://github.com/QwenLM/Qwen-Image-2.1
S4	Alibaba Qwen. <i>Qwen-Image-2.1 model card</i> . Hugging Face, accessed 22 September 2026. https://huggingface.co/Qwen/Qwen-Image-2.1
S5	Stable AI and Tsinghua. <i>LimiX-2 model card</i> . Hugging Face, 16 September 2026. https://huggingface.co/stable-ai/LimiX-2
S6	Zhang, X. et al. <i>LimiX-2: A Contextual Mechanism Network Towards General Structured-Data Intelligence</i> . arXiv:2609.17488, 15 September 2026. https://arxiv.org/abs/2609.17488
S7	Google DeepMind. <i>Introducing Gemini 3.8 Live and 3.8 Live Extended Thinking</i> . Lab blog, 15 September 2026 (updated 17 September). https://blog.google/innovation-and-ai/models-and-research/gemini-models/gemini-3-8-live-gemini-3-8-live-extended-thinking/
S8	Lin, Y. et al. <i>The Other Half of the Memory Wall: Serving 35B MoEs from SSD with Trained Routing Prediction</i> . arXiv:2609.18063, 16 September 2026. https://arxiv.org/abs/2609.18063
S9	Liu, H. et al. <i>SoL-Pi: Recursively Scaling Auto-Research Loops for Efficient Agent Harness</i> . arXiv:2609.20519, 17 September 2026. https://arxiv.org/abs/2609.20519
S10	DeepSeek-AI. <i>DeepSeek-V4.1-Flash: Pushing the Limits of KV Cache Compression</i> . arXiv:2609.19969, 17 September 2026. https://arxiv.org/abs/2609.19969
S11	Yang, Y. et al. <i>When EOS Tokens Disagree: Understanding Length Inflation in On-Policy Distillation</i> . arXiv:2609.20511, 17 September 2026. https://arxiv.org/abs/2609.20511
S12	OpenAI. <i>Our framework for reporting model misalignment</i> . Lab blog, 16 September 2026. https://openai.com/index/model-misalignment-reporting-framework/
S13	Anthropic. <i>Measurements for understanding the pace of AI development inside frontier labs</i> . Anthropic Institute, 17 September 2026. https://www.anthropic.com/institute/measuring-pace-of-ai-development
S14	Anthropic. <i>Partnering with Accenture on embedded evaluation</i> . Lab blog, 18 September 2026. https://www.anthropic.com/news/accenture-embedded-evaluation
S15	SK hynix. <i>No plans have been confirmed regarding the reported talks between SK hynix and Intel on memory chip production in the U.S.</i> Company statement, 16 September 2026. https://news.skhynix.com/en/fact-10/
S16	TechCrunch. <i>SK Hynix reportedly in talks with Intel to build memory chips in US</i> . Press, 16 September 2026. https://techcrunch.com/2026/09/16/sk-hynix-reportedly-in-talks-with-intel-to-build-memory-chips-in-us/

Compiled 22 September 2026. Every figure in this brief was read from the linked primary source on that date. Corrections: Ali Habibullah.